

Various_architectures_on_MNIST

April 1, 2019

```
In [2]: %matplotlib inline
```

```
import warnings
warnings.filterwarnings('ignore')

from keras.utils import np_utils
from keras.datasets import mnist
from keras.initializers import RandomNormal

# Model
from keras.models import Sequential

# Layers
from keras.layers import Dense
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

import seaborn as sns
import matplotlib.pyplot as plt
```

Using TensorFlow backend.

```
In [3]: # MNIST Data Fetching and Preprocessing
```

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()

# Convert 28*28 image into 784 size 1D tensor
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1] * X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1] * X_test.shape[2])

# Convert labels into one-hot encoded vectors

y_train = np_utils.to_categorical(y_train, 10)
y_test = np_utils.to_categorical(y_test, 10)

# Normalize the input data using simple min max normalization
```

```
X_train = X_train / 255
X_test = X_test / 255
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

```
In [0]: def plot_dynamic(x, vy, ty, ax, fig, colors=['b']):
        ax.plot(x, vy, 'r', label='Validation Loss')
        ax.plot(x, ty, 'b', label='Train Loss')
        plt.legend()
        plt.grid()
        fig.canvas.draw()
```

```
In [0]: # Plotting Training/Validation Loss
```

```
def plot_loss(history):
    fig, ax = plt.subplots(1, 1)
    ax.set_xlabel("Epochs")
    ax.set_ylabel("Softmax Cross Entropy Loss")

    x = list(range(1, n_epochs+1))
    vy = history.history['val_loss']
    ty = history.history['loss']

    plot_dynamic(x, vy, ty, ax, fig)
```

```
In [0]: '''
def plot_weight_distribution(weights):
    h1 = weights[0].flatten().reshape(-1, 1)
    h2 = weights[2].flatten().reshape(-1, 1)
    out = weights[4].flatten().reshape(-1, 1)
    fig = plt.figure()
    plt.title("Training Weights Distribution")

    plt.subplot(1, 3, 1)
    plt.title("Trained Weights")
    ax = sns.violinplot(y=h1, color='b')
    plt.xlabel("Layer 1")

    plt.subplot(1, 3, 2)
    plt.title("Trained Weights")
    ax = sns.violinplot(y=h2, color='r')
    plt.xlabel("Layer 2")

    plt.subplot(1, 3, 3)
    plt.title("Trained Weights")
    ax = sns.violinplot(y=out, color='g')
    plt.xlabel('Out Layer')
```

```

        plt.show()
    '''
    # Plot function for weight distribution

    def plot_weight_distribution(weights, hidden_layers=None):
        colors = ['b', 'r', 'g', 'y', 'm']
        fig = plt.figure()
        plt.title("Training Weights Distribution")

        for i in range(0, hidden_layers+1):
            layer_weights = weights[i*2].flatten().reshape(-1, 1)
            plt.subplot(1, hidden_layers+1, i+1)
            plt.title("Trained Weights")
            ax = sns.violinplot(y=layer_weights, color=colors[i%5])
            if i == hidden_layers:
                plt.xlabel("Out Layer")
            else:
                plt.xlabel("Layer {}".format(i+1))
        plt.show()

```

In [0]: # Model Parameters

```

input_dim = 784
output_dim = 10
batch_size = 100
n_epochs = 20

```

MLP (2-hidden layers) + Adam Optimizer + ReLU activations

In [12]: # Model Architecture

```

# Hidden layer 1 --> 256

# Hidden layer 2 --> 64

# Since we are using ReLU activation,
# we will use He-initialization.

model = Sequential()

# Hidden Layer 1
model.add(Dense(256, input_shape=(input_dim,), activation='relu',
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.055,
                                                    seed=None)))

# Hidden Layer 2

```

```

model.add(Dense(64, activation='relu',
                kernel_initializer=RandomNormal(mean=0.0, stddev=0.088,
                                                seed=None)))

# Output SoftMax Layer
model.add(Dense(output_dim, activation='softmax'))

# Defining optimizer, loss function and evaluation metric
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()

```

```

-----
Layer (type)                 Output Shape              Param #
=====
dense_13 (Dense)             (None, 256)               200960
-----
dense_14 (Dense)             (None, 64)                16448
-----
dense_15 (Dense)             (None, 10)                650
=====
Total params: 218,058
Trainable params: 218,058
Non-trainable params: 0
-----

```

```

In [0]: history = model.fit(X_train, y_train, batch_size=batch_size, epochs=n_epochs,
                             verbose=0, validation_data=(X_test, y_test))

```

```

# Test Loss and Accuracy
score = model.evaluate(X_test, y_test, verbose=0)
print("Test Loss : ", score[0])
print("Test Accuracy : ", score[1])

```

```


Test Loss : 0.10167545924528686
Test Accuracy : 0.9802

```


```

In [0]: plot_loss(history)

```



```
In [0]: weights = model.get_weights()
        plot_weight_distribution(weights)
```



MLP (2-hidden layers) + Adam Optimizer + ReLU activations + BatchNorm + Dropout

In [13]: # Model Architecture

```
# Hidden layer 1 --> 256

# Batch Normalization Layer
# Dropout layer with dropout_rate = 0.3

# Hidden layer 2 --> 64

# Batch Normalization Layer
# Dropout layer with dropout_rate = 0.3

# Since we are using ReLU activation,
# we will use He-initialization.

model = Sequential()

# Hidden Layer 1
model.add(Dense(256, input_shape=(input_dim,), activation='relu',
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.055,
                                                    seed=None)))

model.add(BatchNormalization())
model.add(Dropout(0.3))

# Hidden Layer 2
model.add(Dense(64, activation='relu',
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.088,
                                                    seed=None)))

model.add(BatchNormalization())
model.add(Dropout(0.3))

# Output SoftMax Layer
model.add(Dense(output_dim, activation='softmax'))

# Defining optimizer, loss function and evaluation metric
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()
```

Layer (type)

Output Shape

Param #

```

=====
dense_16 (Dense)                (None, 256)                200960
-----
batch_normalization_11 (Batch Normalization) (None, 256)                1024
-----
dropout_11 (Dropout)            (None, 256)                 0
-----
dense_17 (Dense)                (None, 64)                 16448
-----
batch_normalization_12 (Batch Normalization) (None, 64)                 256
-----
dropout_12 (Dropout)            (None, 64)                 0
-----
dense_18 (Dense)                (None, 10)                 650
=====
Total params: 219,338
Trainable params: 218,698
Non-trainable params: 640
-----

```

```

In [0]: history = model.fit(X_train, y_train, batch_size=batch_size, epochs=n_epochs,
                             verbose=0, validation_data=(X_test, y_test))

```

```

    # Test Loss and Accuracy
    score = model.evaluate(X_test, y_test, verbose=0)
    print("Test Loss : ", score[0])
    print("Test Accuracy : ", score[1])

```

```



Test Loss : 0.06912141052252264
Test Accuracy : 0.9816

```

```

In [0]: plot_loss(history)
        weights = model.get_weights()
        plot_weight_distribution(weights)

```



MLP (3-hidden layers) + Adam Optimizer + ReLU activations


```

In [14]: # Model Architecture

# Hidden layer 1 --> 1024

# Hidden layer 2 --> 512

# Hidden layer 3 --> 256

# Since we are using ReLU activation,
# we will use He-initialization.

model = Sequential()

# Hidden Layer 1
model.add(Dense(1024, input_shape=(input_dim,), activation='relu',
                kernel_initializer=RandomNormal(mean=0.0, stddev=0.050,
                seed=None)))

# Hidden Layer 2
model.add(Dense(512, activation='relu',
                kernel_initializer=RandomNormal(mean=0.0, stddev=0.044,
                seed=None)))

# Hidden Layer 3
model.add(Dense(256, activation='relu',
                kernel_initializer=RandomNormal(mean=0.0, stddev=0.062,
                seed=None)))

# Output SoftMax Layer
model.add(Dense(output_dim, activation='softmax'))

# Defining optimizer, loss function and evaluation metric
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()

```

Layer (type)	Output Shape	Param #
dense_19 (Dense)	(None, 1024)	803840
dense_20 (Dense)	(None, 512)	524800

dense_21 (Dense)	(None, 256)	131328
------------------	-------------	--------

dense_22 (Dense)	(None, 10)	2570
------------------	------------	------

Total params: 1,462,538

Trainable params: 1,462,538

Non-trainable params: 0

```
In [0]: history = model.fit(X_train, y_train, batch_size=batch_size, epochs=n_epochs,  
                           verbose=0, validation_data=(X_test, y_test))
```

```
# Test Loss and Accuracy
```

```
score = model.evaluate(X_test, y_test, verbose=0)
```


```
print("Test Loss : ", score[0])
```

```
print("Test Accuracy : ", score[1])
```

Test Loss : 0.09790483784366265

Test Accuracy : 0.9812

```
In [0]: plot_loss(history)
```



```
In [0]: weights = model.get_weights()  
        plot_weight_distribution(weights, hidden_layers=5)
```

MLP (3-hidden layers) + Adam Optimizer + ReLU activations + BatchNorm + Dropout

In [15]: # Model Architecture

```
# Hidden layer 1 --> 1024

# Batch Normalization Layer
# Dropout layer with dropout_rate = 0.5

# Hidden layer 2 --> 512

# Batch Normalization Layer
# Dropout layer with dropout_rate = 0.5

# Hidden layer 3 --> 256

# Batch Normalization Layer
# Dropout layer with dropout_rate = 0.5

model = Sequential()

# Hidden Layer 1
model.add(Dense(1024, input_shape=(input_dim,), activation='relu',
                kernel_initializer=RandomNormal(mean=0.0, stddev=0.050,
                seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

# Hidden Layer 2
model.add(Dense(512, activation='relu',
                kernel_initializer=RandomNormal(mean=0.0, stddev=0.044,
                seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

# Hidden Layer 3
model.add(Dense(256, activation='relu',
                kernel_initializer=RandomNormal(mean=0.0, stddev=0.062,
                seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

# Output SoftMax Layer
model.add(Dense(output_dim, activation='softmax'))
```

```
# Defining optimizer, loss function and evaluation metric
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])
```

```
model.summary()
```

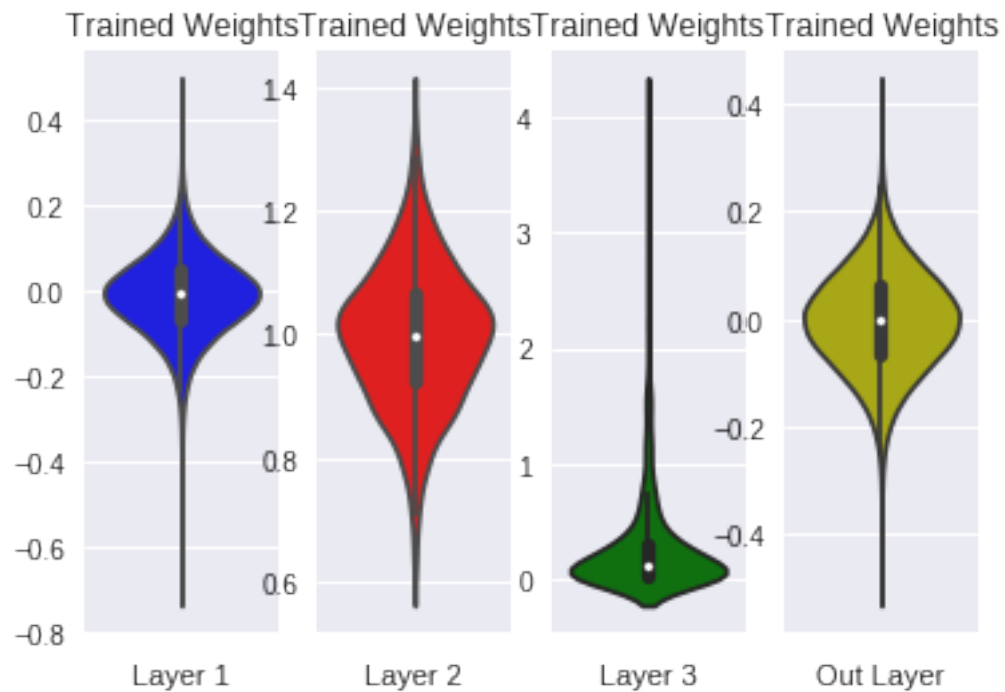

Layer (type)	Output Shape	Param #
dense_23 (Dense)	(None, 1024)	803840
batch_normalization_13 (Batch Normalization)	(None, 1024)	4096
dropout_13 (Dropout)	(None, 1024)	0
dense_24 (Dense)	(None, 512)	524800
batch_normalization_14 (Batch Normalization)	(None, 512)	2048
dropout_14 (Dropout)	(None, 512)	0
dense_25 (Dense)	(None, 256)	131328
batch_normalization_15 (Batch Normalization)	(None, 256)	1024
dropout_15 (Dropout)	(None, 256)	0
dense_26 (Dense)	(None, 10)	2570
Total params: 1,469,706		
Trainable params: 1,466,122		
Non-trainable params: 3,584		

```
In [0]: history = model.fit(X_train, y_train, batch_size=batch_size, epochs=n_epochs,
                           verbose=0, validation_data=(X_test, y_test))
```

```
# Test Loss and Accuracy
score = model.evaluate(X_test, y_test, verbose=0)
print("Test Loss : ", score[0])
print("Test Accuracy : ", score[1])
```

```
Test Loss : 0.05362152259176073
Test Accuracy : 0.985
```

```
In [0]: plot_loss(history)
weights = model.get_weights()
plot_weight_distribution(weights, hidden_layers=3)
```



MLP (5-hidden layers) + Adam Optimizer + ReLU activations

In [16]: *# Model Architecture*

Hidden layer 1 --> 2048

Hidden layer 2 --> 1024

Hidden layer 3 --> 256

Hidden layer 4 --> 128

Hidden layer 5 --> 64

*# Since we are using ReLU activation,
we will use He-initialization.*

model = Sequential()

Hidden Layer 1

model.add(Dense(2048, input_shape=(input_dim,), activation='relu',
kernel_initializer=RandomNormal(mean=0.0, stddev=0.050,
seed=None)))

Hidden Layer 2

model.add(Dense(1024, activation='relu',
kernel_initializer=RandomNormal(mean=0.0, stddev=0.031,
seed=None)))

Hidden Layer 3

model.add(Dense(256, activation='relu',
kernel_initializer=RandomNormal(mean=0.0, stddev=0.044,
seed=None)))

Hidden Layer 4

model.add(Dense(128, activation='relu',
kernel_initializer=RandomNormal(mean=0.0, stddev=0.088,
seed=None)))

Hidden Layer 5

model.add(Dense(64, activation='relu',
kernel_initializer=RandomNormal(mean=0.0, stddev=0.125,

```
seed=None)))
```

```
# Output SoftMax Layer
```

```
model.add(Dense(output_dim, activation='softmax'))
```

```
# Defining optimizer, loss function and evaluation metric
```

```
model.compile(optimizer='adam', loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

```
model.summary()
```

```
-----  
Layer (type)                 Output Shape              Param #  
-----  
dense_27 (Dense)             (None, 2048)              1607680  
-----  
dense_28 (Dense)             (None, 1024)              2098176  
-----  
dense_29 (Dense)             (None, 256)               262400  
-----  
dense_30 (Dense)             (None, 128)               32896  
-----  
dense_31 (Dense)             (None, 64)                8256  
-----  
dense_32 (Dense)             (None, 10)                650  
-----  
Total params: 4,010,058  
Trainable params: 4,010,058  
Non-trainable params: 0  
-----
```

```
In [0]: history = model.fit(X_train, y_train, batch_size=batch_size, epochs=n_epochs,  
                           verbose=0, validation_data=(X_test, y_test))
```

```
# Test Loss and Accuracy
```

```
score = model.evaluate(X_test, y_test, verbose=0)
```

```
print("Test Loss : ", score[0])
```

```
print("Test Accuracy : ", score[1])
```


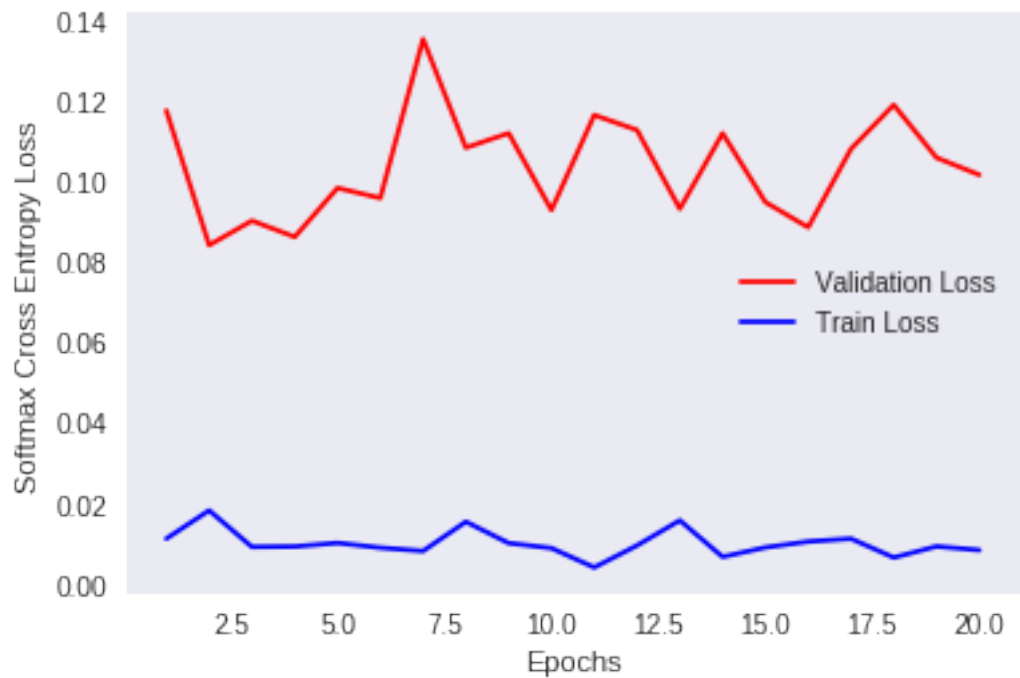
```
Test Loss : 0.10174559558869126
```

```
Test Accuracy : 0.9858
```

```
In [0]: plot_loss(history)
```

```
weights = model.get_weights()
```

```
plot_weight_distribution(weights, hidden_layers=5)
```



MLP (5-hidden layers) + Adam Optimizer + ReLU activations + BatchNorm + Dropout


```

In [17]: # Model Architecture

# Hidden layer 1 --> 2048

# Batch Normalization Layer
# Dropout layer with dropout_rate = 0.5

# Hidden layer 2 --> 1024

# Batch Normalization Layer
# Dropout layer with dropout_rate = 0.5

# Hidden layer 3 --> 256

# Batch Normalization Layer
# Dropout layer with dropout_rate = 0.5

# Hidden layer 4 --> 128

# Batch Normalization Layer
# Dropout layer with dropout_rate = 0.5

# Hidden layer 5 --> 64

# Batch Normalization Layer
# Dropout layer with dropout_rate = 0.5

# Since we are using ReLU activation,
# we will use He-initialization.

model = Sequential()

# Hidden Layer 1
model.add(Dense(2048, input_shape=(input_dim,), activation='relu',
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.050,
                                                    seed=None)))

model.add(BatchNormalization())
model.add(Dropout(0.5))

# Hidden Layer 2
model.add(Dense(1024, activation='relu',
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.031,
                                                    seed=None)))

model.add(BatchNormalization())
model.add(Dropout(0.5))

```

```

# Hidden Layer 3
model.add(Dense(256, activation='relu',
                kernel_initializer=RandomNormal(mean=0.0, stddev=0.044,
                                                seed=None)))

model.add(BatchNormalization())
model.add(Dropout(0.5))

# Hidden Layer 4
model.add(Dense(128, activation='relu',
                kernel_initializer=RandomNormal(mean=0.0, stddev=0.088,
                                                seed=None)))

model.add(BatchNormalization())
model.add(Dropout(0.5))

# Hidden Layer 5
model.add(Dense(64, activation='relu',
                kernel_initializer=RandomNormal(mean=0.0, stddev=0.125,
                                                seed=None)))

model.add(BatchNormalization())
model.add(Dropout(0.5))

# Output SoftMax Layer
model.add(Dense(output_dim, activation='softmax'))

# Defining optimizer, loss function and evaluation metric
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()

```

Layer (type)	Output Shape	Param #
dense_33 (Dense)	(None, 2048)	1607680
batch_normalization_16 (Batch Normalization)	(None, 2048)	8192
dropout_16 (Dropout)	(None, 2048)	0
dense_34 (Dense)	(None, 1024)	2098176
batch_normalization_17 (Batch Normalization)	(None, 1024)	4096
dropout_17 (Dropout)	(None, 1024)	0

dense_35 (Dense)	(None, 256)	262400
batch_normalization_18 (Batch Normalization)	(None, 256)	1024
dropout_18 (Dropout)	(None, 256)	0
dense_36 (Dense)	(None, 128)	32896
batch_normalization_19 (Batch Normalization)	(None, 128)	512
dropout_19 (Dropout)	(None, 128)	0
dense_37 (Dense)	(None, 64)	8256
batch_normalization_20 (Batch Normalization)	(None, 64)	256
dropout_20 (Dropout)	(None, 64)	0
dense_38 (Dense)	(None, 10)	650

=====
 Total params: 4,024,138
 Trainable params: 4,017,098
 Non-trainable params: 7,040
 =====

```
In [11]: history = model.fit(X_train, y_train, batch_size=batch_size, epochs=n_epochs,
                             verbose=0, validation_data=(X_test, y_test))
```

```

# Test Loss and Accuracy
score = model.evaluate(X_test, y_test, verbose=0)
print("Test Loss : ", score[0])
print("Test Accuracy : ", score[1])



plot_loss(history)
weights = model.get_weights()
plot_weight_distribution(weights, hidden_layers=5)

```

```

Test Loss : 0.05905918826502748
Test Accuracy : 0.9854

```



1 Observations

1. 2 hidden layer(256-64) --> Loss = 0.10, Accuracy = 98.02%
2. 2 hidden layer(256-64) --> Loss = 0.069, Accuracy = 98.16%
3. 3 hidden layer(1024-512-256) --> Loss = 0.097, Accuracy = 98.12%
4. 3 hidden layer(1024-512-256) --> Loss = 0.053, Accuracy = 98.50%
5. 5 hidden layer(2048-1024-256-128-64) --> Loss = 0.10, Accuracy = 98.58%
6. 5 hidden layer(2048-1024-256-128-64) --> Loss = 0.059, Accuracy = 98.54%

Getting best accuracy with 5 hidden layers and 3 hidden layers with batch normalization and dropout rate of 0.5.

Also the Train and Test Loss follow each other very closely when we use dropout and batch normalization layers.